

Analysis of QIIME2 output

Lev Litichevskiy

January 12, 2023

The goal of this notebook is to make several key plots based on the output of QIIME2, specifically the counts matrix and the taxonomy annotations. QIIME2 has plug-ins for doing all of these analyses, but I prefer to do them in R.

The plots that we will make are:

- 1) PCoA
- 2) Barplots at different taxonomic levels
- 3) Alpha diversity
- 4) TODO: identifying taxons different between two groups with ANCOM or DESeq2

Search the document for “CHANGEME” to see which lines of code will require project-specific tweaks.

This demo dataset contains gut microbiome samples from young, middle-aged, and old mice before and after undergoing fecal microbial transfer (FMT) of a young microbiome. Before FMT, we expect to see samples separate by age, but after FMT, the samples look quite similar to each other because they received the same FMT input.

Load libraries

```
library(tidyverse)
library(cowplot)
library(ggpubr)
library(phyloseq)
library(speedyseq) # to speed up the tax_glom function
library(vegan)

# use black-and-white theme, and increase default font size
theme_set(theme_bw(base_size=15))
```

Import metadata

This file should contain any important metadata about your samples of interest.

```
(meta.df <- read.table("data/demo_qiime2_metadata.tsv", sep="\t", header=T))
```

```
##      Sample.ID      Age Timepoint
## 1          S1 Middle   Pre-FMT
## 2          S2 Middle   Post-FMT
## 3          S3 Middle   Pre-FMT
## 4          S4 Middle   Post-FMT
## 5          S5 Middle   Pre-FMT
```

```
## 6      S6 Middle Post-FMT
## 7      S7 Middle  Pre-FMT
## 8      S8 Middle Post-FMT
## 9      S9      Old  Pre-FMT
## 10     S10     Old Post-FMT
## 11     S11     Old  Pre-FMT
## 12     S12     Old Post-FMT
## 13     S13     Old  Pre-FMT
## 14     S14     Old  Pre-FMT
## 15     S15     Old Post-FMT
## 16     S16 Young  Pre-FMT
## 17     S17 Young Post-FMT
## 18     S18 Young  Pre-FMT
## 19     S19 Young Post-FMT
## 20     S20 Young  Pre-FMT
## 21     S21 Young Post-FMT
## 22     S22 Young  Pre-FMT
## 23     S23 Young  Pre-FMT
## 24     S24 Young Post-FMT
## 25     S25 Young  Pre-FMT
```

- The sample identifier should be in a column called `Sample.ID`
- If it's not, you can search and replace `Sample.ID` with whatever else (e.g. `sample_id`)

Import taxonomy

```
tax.df <- read.table("data/demo_qiime2_taxonomy.tsv", sep="\t", header=T)

# create new column for each taxonomy level
tax.df <- tax.df %>%
  separate(Taxon, c("kingdom", "phylum", "class", "order", "family", "genus", "species"),
    sep=";", [[alpha:]]__, remove=T, fill="right")

# take a peek at the taxonomy
tax.df[1:3, ]
```

```
##               Feature.ID      kingdom      phylum      class
## 1 9c9bfd420fffa8772f8982255054c692 d__Bacteria Bacteroidota Bacteroidia
## 2 1cb85e10fa890fda99460d5703c18abd d__Bacteria  Firmicutes    Bacilli
## 3 abcda143e8418f4e2b05862424b1a1a2 d__Bacteria Bacteroidota Bacteroidia
##               order      family      genus      species
## 1  Bacteroidales  Bacteroidaceae  Bacteroides Bacteroides_thetaiotaomicron
## 2  Lactobacillales Lactobacillaceae Lactobacillus Lactobacillus_johnsonii
## 3  Bacteroidales  Muribaculaceae Muribaculaceae uncultured_bacterium
## Confidence
## 1 0.9987711
## 2 0.9943885
## 3 0.7073147
```

Import counts

In the default output of QIIME2, the second line of the counts table starts with `#` before `OTU.ID`, which causes R to treat this line as a comment. You will need to delete the `#` at the start of the second line before

trying to import into R.

```
counts.df <- read.table("data/demo_qiime2_feature_table.tsv", sep="\t", header=T)

# set the taxon names as the rownames
counts.df <- column_to_rownames(counts.df, var="OTU.ID")

# take a peek at the counts
counts.df[1:5, 1:5]
```

```
##                               S1      S2      S3      S4      S5
## 9c9bfd420fffa8772f8982255054c692  114  6687  118  7437    50
## 1cb85e10fa890fda99460d5703c18abd 4994 43562 5507 84937 10130
## abcda143e8418f4e2b05862424b1a1a2 4998 10860 5338 42894  2206
## 1bb73fe7e3fdfcc431d59cc403af855f  402    17   870    34  2576
## a2213eede305d5a387ec2a9e56412657    9  6330   11   112     0
```

```
dim(counts.df)
```

```
## [1] 2593    25
```

We have 2593 taxons x 25 samples.

Confirm that all samples are in the metadata

```
all(colnames(counts.df) %in% meta.df$Sample.ID)
```

```
## [1] TRUE
```

(optional) Filtration

Sometimes, we need to discard samples that received too few reads. Or we want to discard taxons that we detected very rarely because we think they're noise or contamination. This section can help you do that.

For example, say that we want to keep samples with at least 10k total counts and taxons detected in at least 5% of samples. These thresholds are arbitrary: change them as necessary.

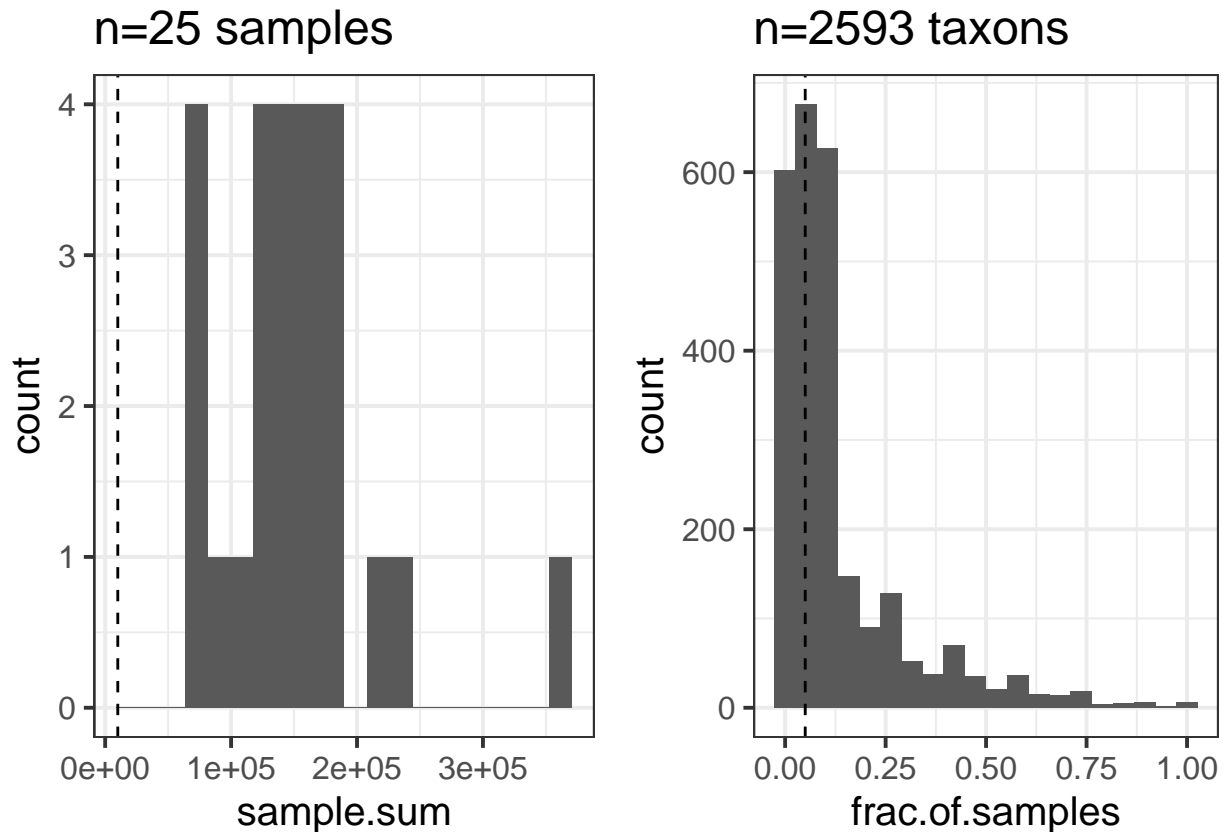
Visualize filtration thresholds

```
# CHANGE ME if you want to change filtration thresholds
SAMPLE.SUM.THRESHOLD <- 10000
TAXON.FRAC.THRESHOLD <- 0.05

p1 <- data.frame(sample.sum=colSums(counts.df)) %>%
  ggplot(aes(sample.sum)) +
  geom_histogram(bins=20) +
  geom_vline(xintercept=SAMPLE.SUM.THRESHOLD, lty=2) +
  labs(title=sprintf("n=%i samples", ncol(counts.df)))

p2 <- data.frame(frac.of.samples=rowSums(counts.df > 0)/ncol(counts.df)) %>%
  ggplot(aes(frac.of.samples)) +
  geom_histogram(bins=20) +
  geom_vline(xintercept=TAXON.FRAC.THRESHOLD, lty=2) +
  labs(title=sprintf("n=%i taxons", nrow(counts.df)))
```

```
plot_grid(p1, p2, nrow=1)
```



Based on these thresholds, we'll keep all samples and discard lots of taxa.

Do filtering

```
samples.to.keep <- colnames(counts.df)[colSums(counts.df) > SAMPLE.SUM.THRESHOLD]
taxons.to.keep <- rownames(counts.df)[rowSums(counts.df > 0)/ncol(counts.df) > TAXON.FRAC.THRESHOLD]
counts.filt.df <- counts.df[taxons.to.keep, samples.to.keep]
```

```
dim(counts.df)
```

```
## [1] 2593 25
```

```
dim(counts.filt.df)
```

```
## [1] 1315 25
```

With these thresholds, we go from 2593 to 1315 taxons, and we don't discard any samples.

I will use the **unfiltered** counts dataframe for the rest of this notebook. If you want to use the filtered data, you'll need to tweak the metadata and taxonomy dataframes (see next code chunk).

Create phyloseq object

We combine all our data into a phyloseq object because the phyloseq package has a number of useful functions that we want to use.

```

# if using the filtered dataframe
# tmp.physeq.meta.df <- meta.df %>% column_to_rownames("Sample.ID")
# physeq.meta.df <- tmp.physeq.meta.df[colnames(counts.filt.df), ]
# tmp.physeq.tax.df <- tax.df %>% column_to_rownames("Feature.ID")
# physeq.tax.df <- tmp.physeq.tax.df[rownames(counts.filt.df), ]

# if using the unfiltered dataframe
physeq.meta.df <- meta.df %>% column_to_rownames("Sample.ID")
physeq.tax.df <- tax.df %>% column_to_rownames("Feature.ID")

physeq <- phyloseq(
  counts.df %>% as.matrix() %>% otu_table(taxa_are_rows=T), # can substitute counts.filt.df here
  physeq.meta.df %>% sample_data(),
  physeq.tax.df %>% as.matrix() %>% tax_table()
)

physeq

## phyloseq-class experiment-level object
## otu_table() OTU Table: [ 2593 taxa and 25 samples ]:
## sample_data() Sample Data: [ 25 samples by 2 sample variables ]:
## tax_table() Taxonomy Table: [ 2593 taxa by 8 taxonomic ranks ]:
## taxa are rows

```

Aggregate to different taxonomy levels

```

physeq.phylum <- physeq %>% tax_glom(taxrank="phylum")
physeq.class <- physeq %>% tax_glom(taxrank="class")
physeq.family <- physeq %>% tax_glom(taxrank="family")
physeq.genus <- physeq %>% tax_glom(taxrank="genus")
physeq.species <- physeq %>% tax_glom(taxrank="species")

```

```
ntaxa(physeq.phylum)
```

```
## [1] 10
```

```
ntaxa(physeq.class)
```

```
## [1] 15
```

```
ntaxa(physeq.family)
```

```
## [1] 70
```

```
ntaxa(physeq.genus)
```

```
## [1] 146
```

```
ntaxa(physeq.species)
```

```
## [1] 261
```


I like to use genus-level data.

PCoA

```
pcoa.genus.res <- physeq.genus %>%  
  
  # compute relative abundance  
  transform_sample_counts(function(OTU) OTU/sum(OTU)) %>%  
  
  # perform PCoA using Bray-Curtis distance  
  ordinate(method="MDS", distance="bray")
```

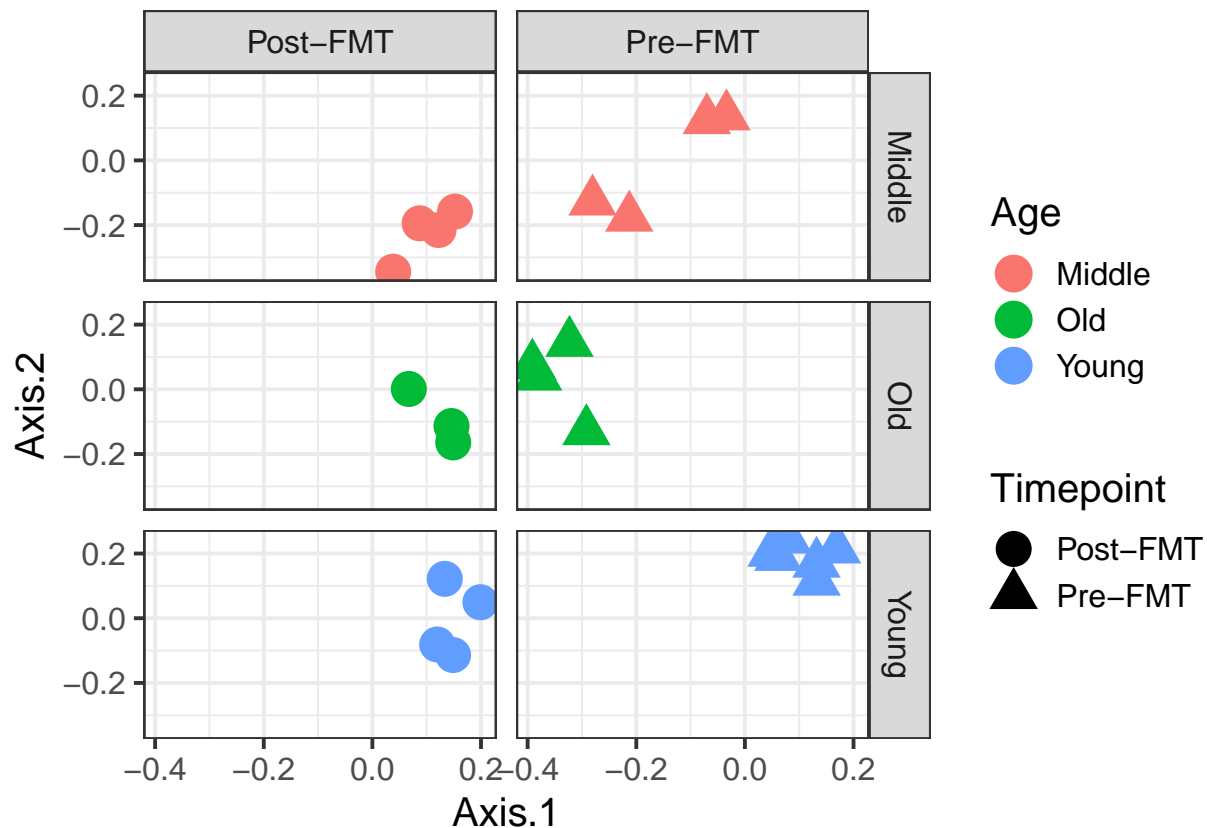
We can use the `plot_ordination` function to make a plot.

```
# CHANGE: adjust color and shape as desired  
plot_ordination(physeq.genus, pcoa.genus.res, type="samples", color="Age", shape="Timepoint") +  
  geom_point(size=6) +  
  facet_wrap(~Timepoint)
```



We see good separation by age before FMT, but not as much separation after FMT. We can also directly get the PCoA coordinates and make a custom plot ourselves.

```
# get PCoA coordinates  
pcoa.genus.df <- plot_ordination(physeq.genus, pcoa.genus.res, type="samples", justDF=T)  
  
# CHANGE: change color, shape, and facet to your variables of interest  
pcoa.genus.df %>%  
  ggplot(aes(x=Axis.1, y=Axis.2, color=Age, shape=Timepoint)) +  
  geom_point(size=6) +  
  facet_grid(Age~Timepoint)
```



Barplots

Phylum

```
# identify the top 4 most abundant phyla
# this is especially important for lower taxonomic levels that
# have too many groups to show in one plot
top.n.phyla <- data.frame(tax_table(physeq.phylum)) [
  names(sort(taxa_sums(physeq.phylum), decreasing=T))[1:4], "phylum"]

physeq.phylum %>%

  # convert to df
  psmelt %>%

  # aggregate less common phyla into "Other"
  mutate(agg.phylum=fct_relevel(
    ifelse(phylum %in% top.n.phyla, phylum, "Other"), "Other", after=Inf)) %>%

  # compute sum of counts per sample
  group_by(Sample) %>%
  mutate(total.abund=sum(Abundance)) %>%

  # compute sum of counts per sample-phylum group
  group_by(Sample, agg.phylum, total.abund) %>%
  summarise(abund=sum(Abundance), .groups="drop") %>%
```

```

# compute relative abundance
mutate(rel.abund=abund/total.abund) %>%

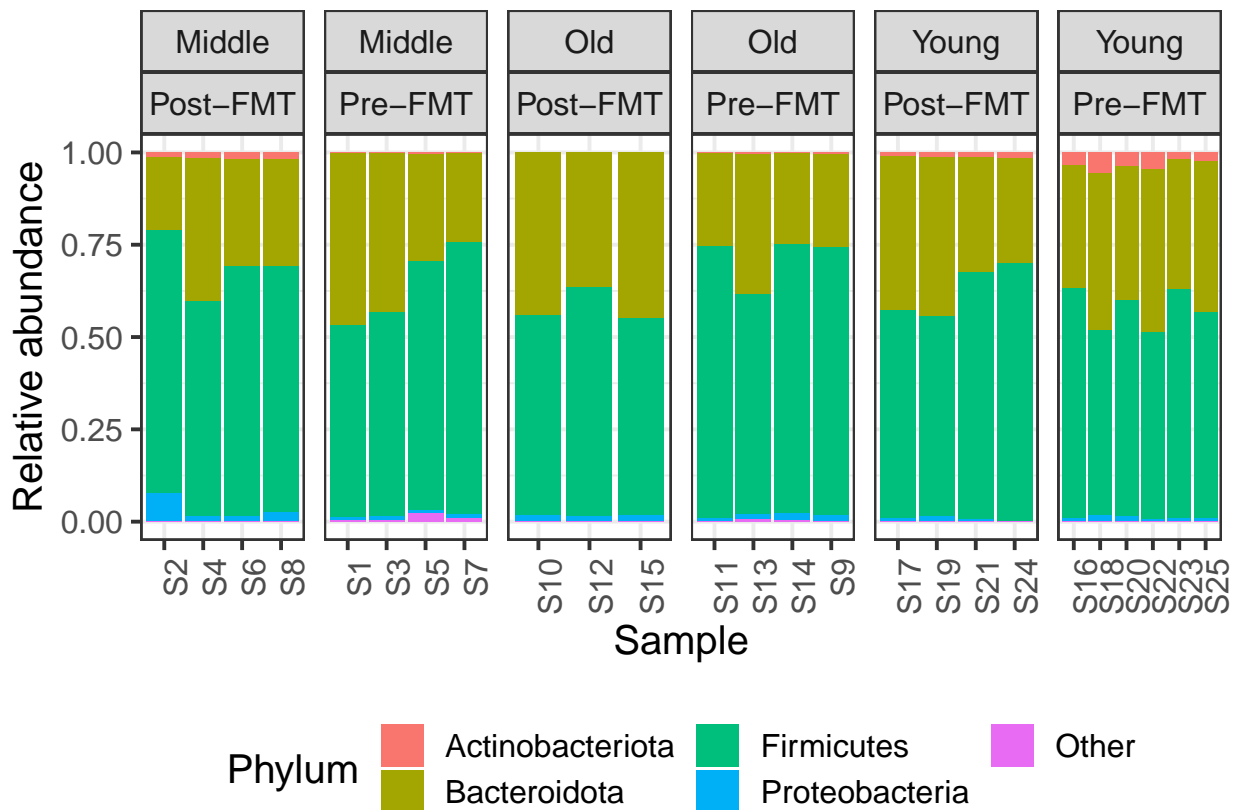
# add back metadata
merge(data.frame(sample_data(physeq.phylum)), by.x="Sample", by.y="row.names") %>%

# plot
ggplot(aes(x=Sample, y=rel.abund, fill=agg.phylum)) +
  geom_bar(stat="identity") +
  labs(y="Relative abundance", fill="Phylum") +

# put legend below plots and rotate x-labels
theme(legend.position="bottom",
      axis.text.x=element_text(angle=90, hjust=1)) +
  guides(fill=guide_legend(nrow=2)) +

# CHANGE ME: modify as necessary
facet_wrap(~Age+Timepoint, nrow=1, scales="free_x")

```



Genus

```

# identify the 9 most abundant genera
top.n.genera <- data.frame(tax_table(physeq.genus)) [
  names(sort(taxa_sums(physeq.genus), decreasing=T))[1:9], "genus"]

physeq.genus %>%

```



```

# convert to df
psmelt %>%

# aggregate less common genera into "Other"
mutate(agg.genus=fct_relevel(
  ifelse(genus %in% top.n.genera, genus, "Other"), "Other", after=Inf)) %>%

# compute sum of counts per sample
group_by(Sample) %>%
mutate(total.abund=sum(Abundance)) %>%

# compute sum of counts per sample-genus group
group_by(Sample, agg.genus, total.abund) %>%
summarise(abund=sum(Abundance), .groups="drop") %>%

# compute relative abundance
mutate(rel.abund=abund/total.abund) %>%


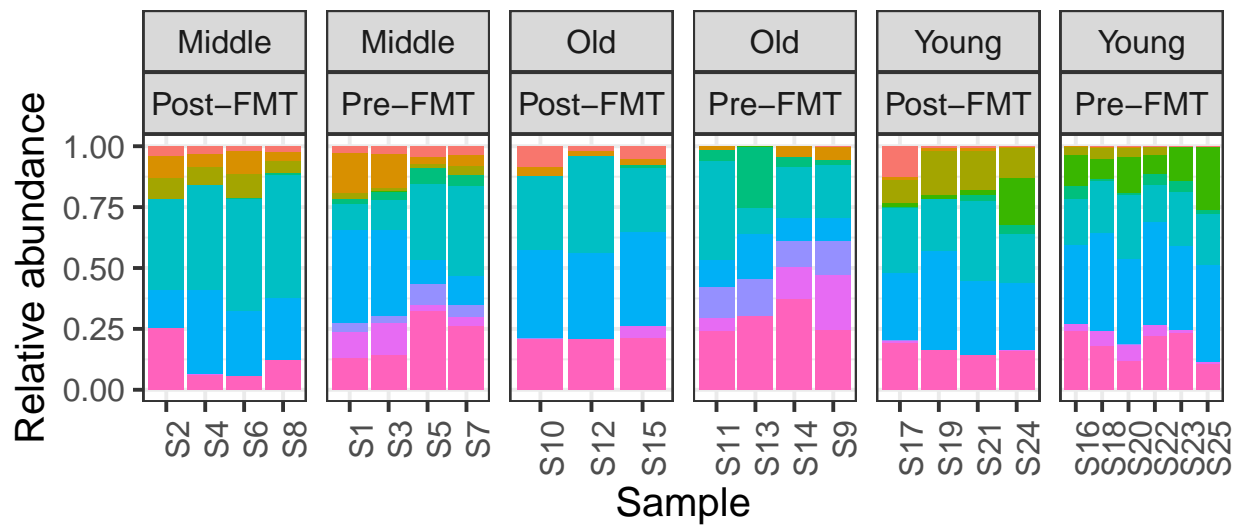
# add back metadata
merge(data.frame(sample_data(physeq.genus)), by.x="Sample", by.y="row.names") %>%

# plot
ggplot(aes(x=Sample, y=rel.abund, fill=agg.genus)) +
geom_bar(stat="identity") +
labs(y="Relative abundance", fill="Genus") +

# put legend below plots and rotate x-labels
theme(legend.position="bottom",
  axis.text.x=element_text(angle=90, hjust=1)) +
guides(fill=guide_legend(nrow=5)) +

# CHANGE ME: modify as necessary
facet_wrap(~Age+Timepoint, nrow=1, scales="free_x")

```



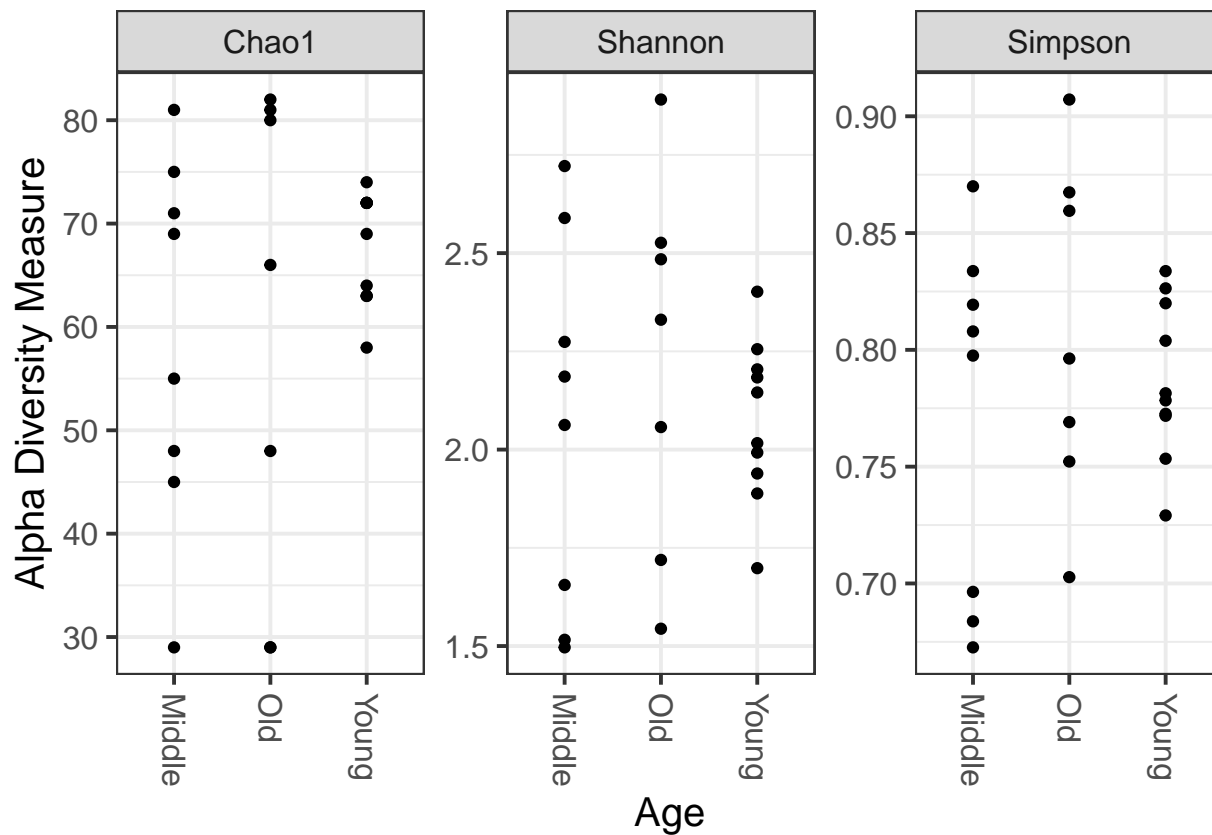
Alpha diversity

Genus

We can use the `plot_richness` command from `phyloseq`.

```
# CHANGE: change x to be your variable of interest
plot_richness(physeq.genus, x="Age", measures=c("Chao1", "Shannon", "Simpson"))
```

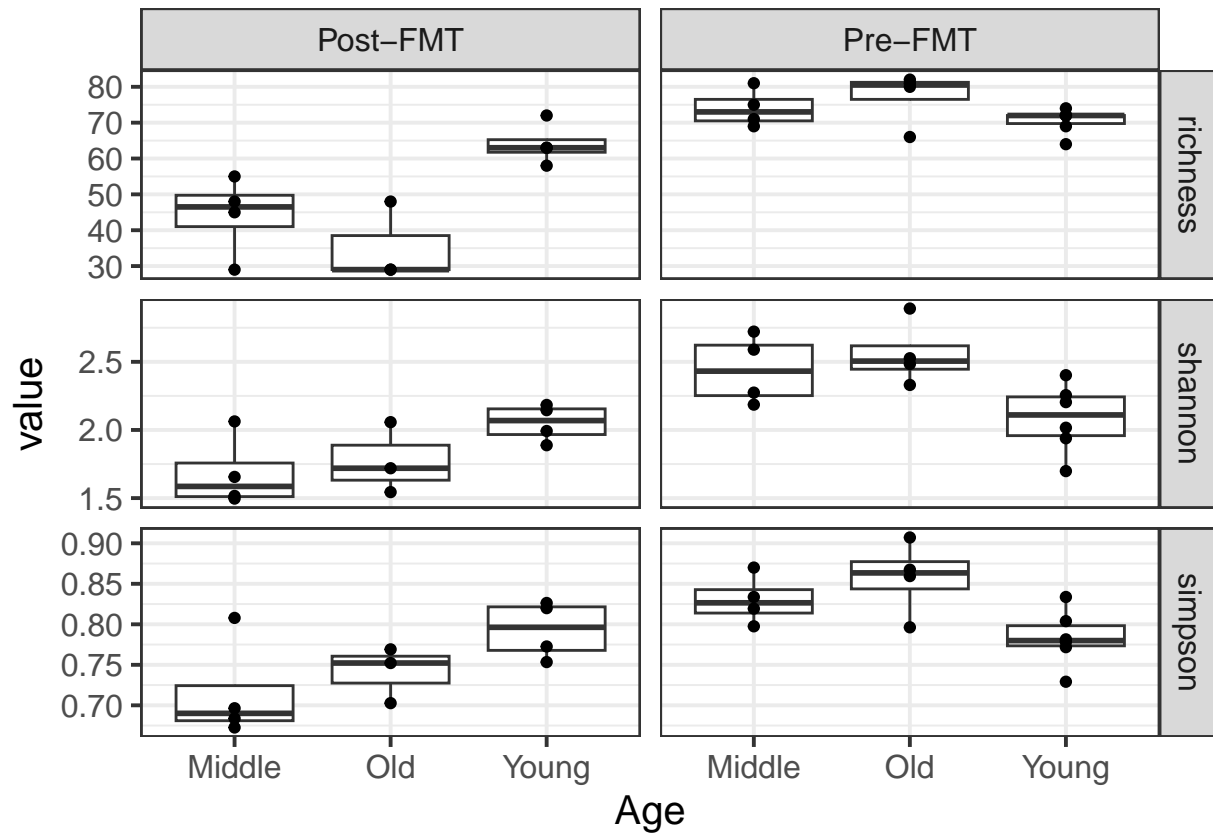
```
## Warning in estimate_richness(physeq, split = TRUE, measures = measures): The data you have provided contains
## any singletons. This is highly suspicious. Results of richness
## estimates (for example) are probably unreliable, or wrong, if you have already
## trimmed low-abundance taxa from the data.
##
## We recommended that you find the un-trimmed data and retry.
```



Or compute the diversity metrics manually.

```
alpha.div.genus.df <- data.frame(
  shannon=diversity(otu_table(physeq.genus), index="shannon", MARGIN=2),
  simpson=diversity(otu_table(physeq.genus), index="simpson", MARGIN=2),
  richness=colSums(otu_table(physeq.genus) != 0)) %>%
  rownames_to_column("Sample.ID") %>%
  pivot_longer(c(shannon, simpson, richness), names_to="diversity", values_to="value") %>%
  merge(meta.df, by="Sample.ID")
```

```
# CHANGE ME: adjust plot as desired
alpha.div.genus.df %>%
  ggplot(aes(x=Age, y=value)) +
  geom_boxplot(outlier.shape=NA) +
  geom_point() +
  facet_grid(diversity ~ Timepoint, scales="free_y")
```



TODO: Use ANCOM or DESeq2 to identify taxons different between two groups

I will get to this some day. The phyloseq website has a decent tutorial about this: <https://joey711.github.io/phyloseq-extensions/DESeq2.html>.